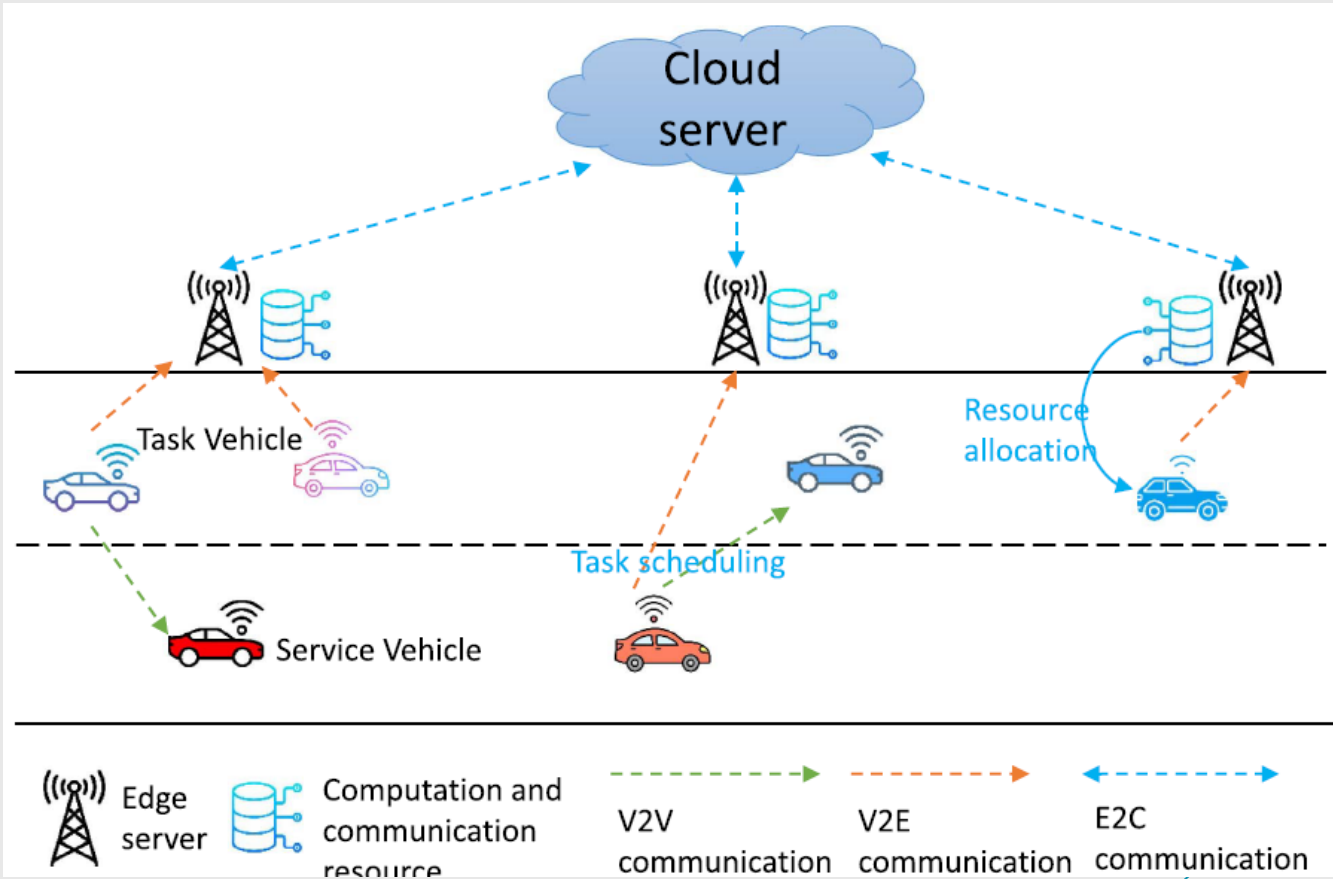




# **EPtask:** Deep Reinforcement Learning Based Energy-Efficient and Priority-Aware Task Scheduling for Dynamic Vehicular Edge Computing

**Mahan Bayhaghi**

# VEC



## EPTask on VEC: PPO vs. DDPG/SAC & Baselines

- **Goal:** Priority-aware, energy-efficient task scheduling in **Vehicular Edge Computing (VEC)**
- **What we built:** A paper-aligned simulator + RL scheduler (PPO) with suggested baselines (SAC, DDPG, heuristic, local-only, offloading-only, random)
- **Key metrics:** deadline miss rate, completion time, mean energy, throughput.
- **Takeaway:** PPO learns smarter offloading & prioritization under mobility, deadlines, and energy limits.

# Problem & Project Objectives

- **Challenge:** Where to compute each task (local/V2V/edge/cloud) and at what priority, under deadlines, mobility, and limited bandwidth/energy
- **Project objectives:**
  - Implement EPTask-style environment
  - Train PPO and compare against DDPG, SAC, heuristic, local-only, offloading-only, random
  - Report miss rate, completion time, energy, throughput
  - show scalability vs. #vehicles
- **Design constraints (per proposal & our choices):**
  - Non-preemptive + EDF
  - Distance-based TX power by default (learned-power kept for ablations).

# Paper Alignment & Assumptions

## •Table III mapping (core params):

•**Radio:** V2I 100 MHz; noise  $\approx -95$  dBm (from  $-174$  dBm/Hz over 100 MHz); vehicle TX  $\approx 1$  dBm (bins for tiers).

•**Compute:** vehicle  $\approx 1$  GHz, edge  $\approx 2$  GHz, cloud  $\approx 5$  GHz (MIPS equivalents).

•**Tasks:** size 2–20 Mb, cycles  $2-20 \times 10^9$ , 4 priority levels; **Poisson** arrivals (bursty).

•**Mobility & links:** 1-D movement; distance  $\rightarrow$  SNR  $\rightarrow$  rate

•**Scheduling rule:** EDF

•**Energy model:** TX + compute energy per task

TABLE III  
EXPERIMENTAL PARAMETERS

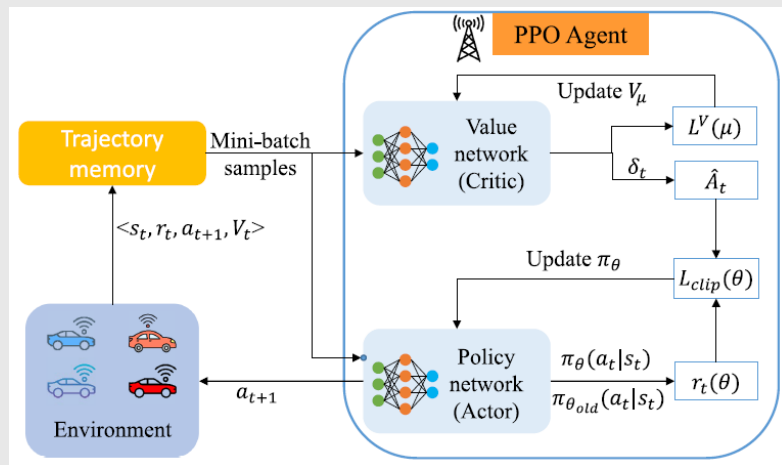
Parameters	Value
Number of vehicles	10-50
Number of edge servers	8
Number of tasks	0-10
Required CPU cycles of the tasks $C$	2-20 cpu cycles
Computation power of edge servers	2 cpu cycles/s
Computation power of vehicles	1 cpu cycles/s
data size $S$	2-20 Mb
Bandwidth of edge server	100 MHz
Speed of vehicles	$\approx 25$ m/s
Transmit power of vehicles	1 dBm
Execution power of vehicles	3-4 dBm
White Gaussian noise	-174 dBm/Hz
Power consumption coefficient $\xi$	$10^{-11}$
Power consumption coefficient $\gamma$	2
Received Signal Strength Indicator $RSSI$	-65 dBm
Transmit antenna gain $TX$	20 dBi
Receive antenna gain $RX$	-8 dBi
Signal attenuation $SA$	7 dB
Working frequency $f$	5GHz
Learning rate	0.0003
Size of Mini-batch	32
Number of steps in each episode	2048
Entropy loss coefficient	0.01

# Simulator Architecture (What's in Code)

- **Env core (env/vec\_env.py):**
  - Builds **Dict observation** (global/task/target/vehicle tensors)
  - **Action** per step on top-K tasks: (*offload target, priority*)
  - **Task generation:** Poisson arrivals; sample\_tasks() from env/generators.py
  - **Link budget:** distance→SNR→rate
  - **Energy/time:** env/models.py (tx/compute time + energy functions)
  - **Reward:**  $-(z\text{-time} + z\text{-energy}) - \lambda \cdot \text{misses}$
- **Metrics & tools:**
  - env/metrics.py, scripts/eval\_metrics.py, scripts/plot\_metrics.py.
  - Grid & scalability: scripts/run\_grid.py, scripts/plot\_vs\_vehicles.py.

# RL Formulation & Methods

- **Action (per top-K tasks):**
  - **Offload target**  $\in \{\text{local, vehicles, edges, cloud}\}$
  - **Priority level**  $\in \{1..4\}$
  - *(default is distance-based, no power head)*
- **Reward:**
  - $(\lambda_1 \cdot \mathbf{z}(\text{completion time}) + \lambda_2 \cdot \mathbf{z}(\text{energy})) - \lambda_3 \cdot \# \text{deadline\_misses}$  (per step)
- **Training setup:** horizon 600; n\_steps=2048, lr=3e-4; eval on 5 episodes;

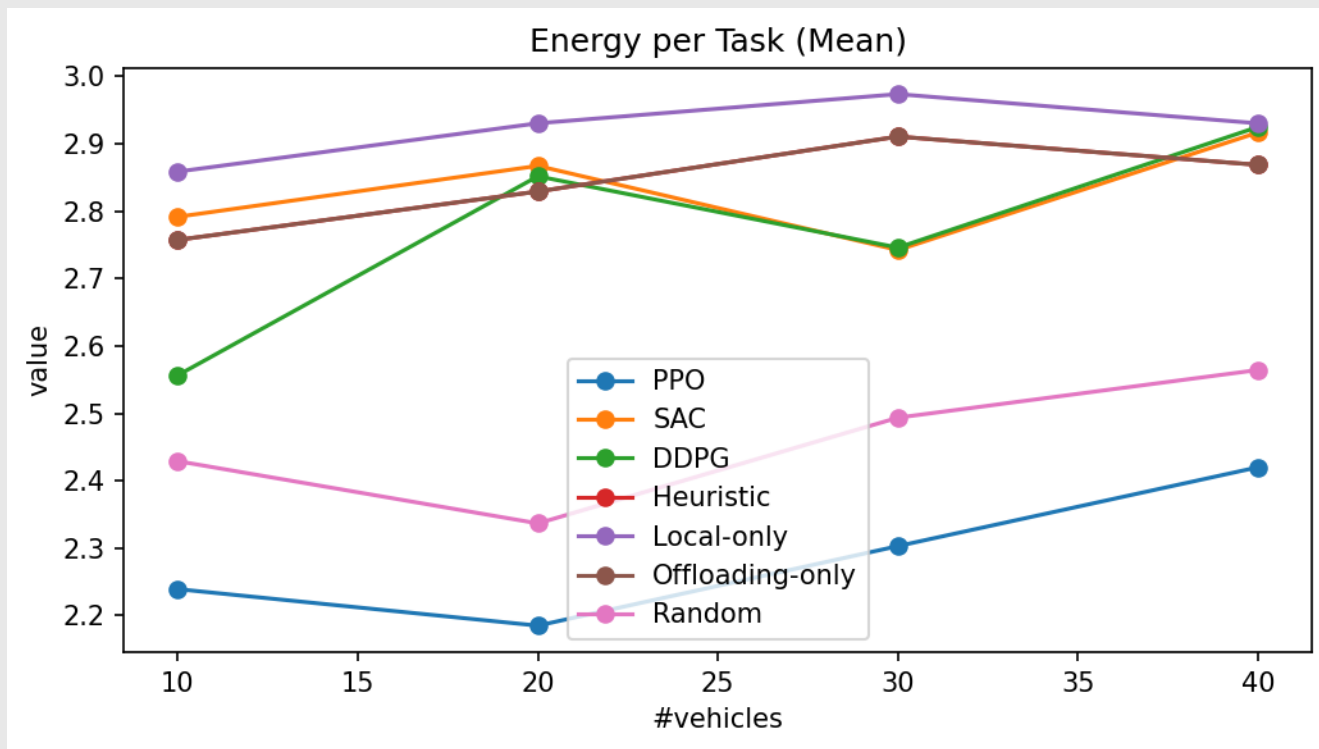


# Results

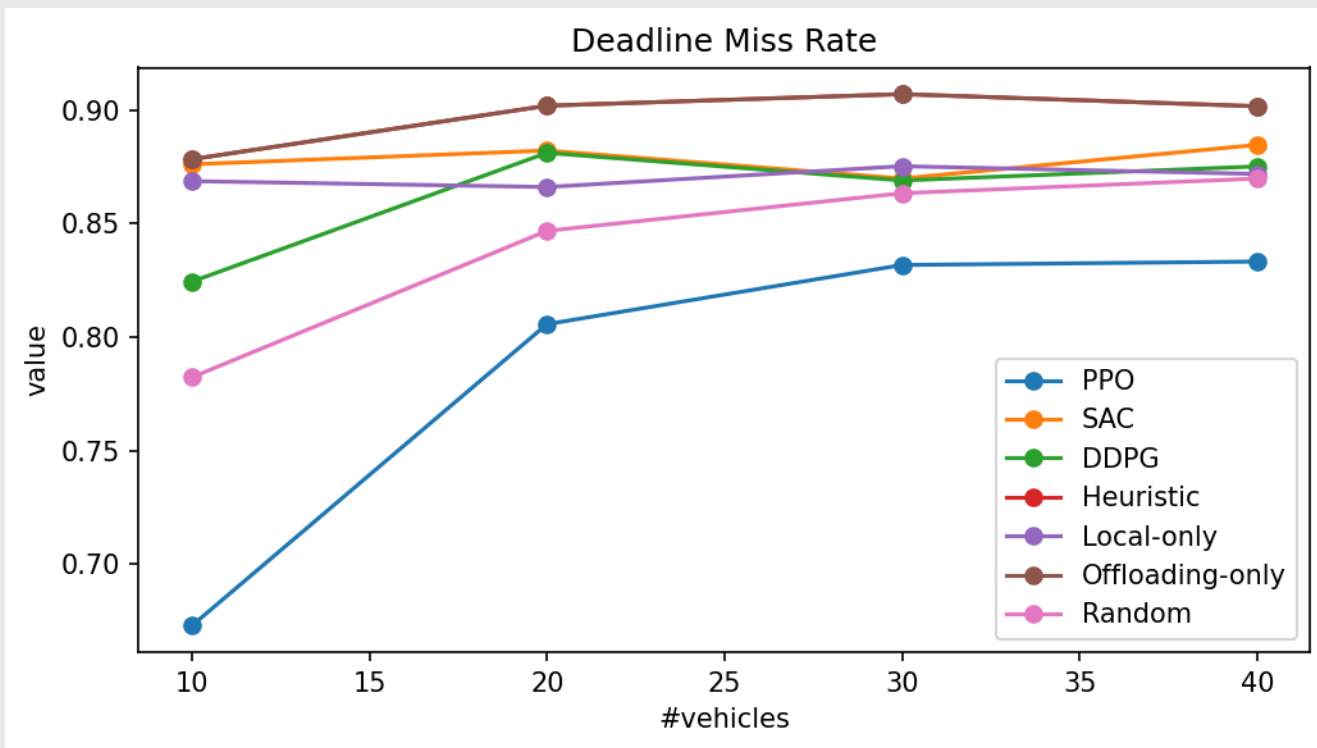
- **Setup:** configs/default.yaml, horizon 600, Poisson arrivals, distance-based power
- **Key findings (qualitative):**
  - **PPO** achieves the **lowest deadline miss rate**
  - Energy per task stays competitive with SAC/DDPG and lower than fixed baselines at load
- **Why:** PPO's discrete offload + priority decisions adapt to bursty queues and mobility.



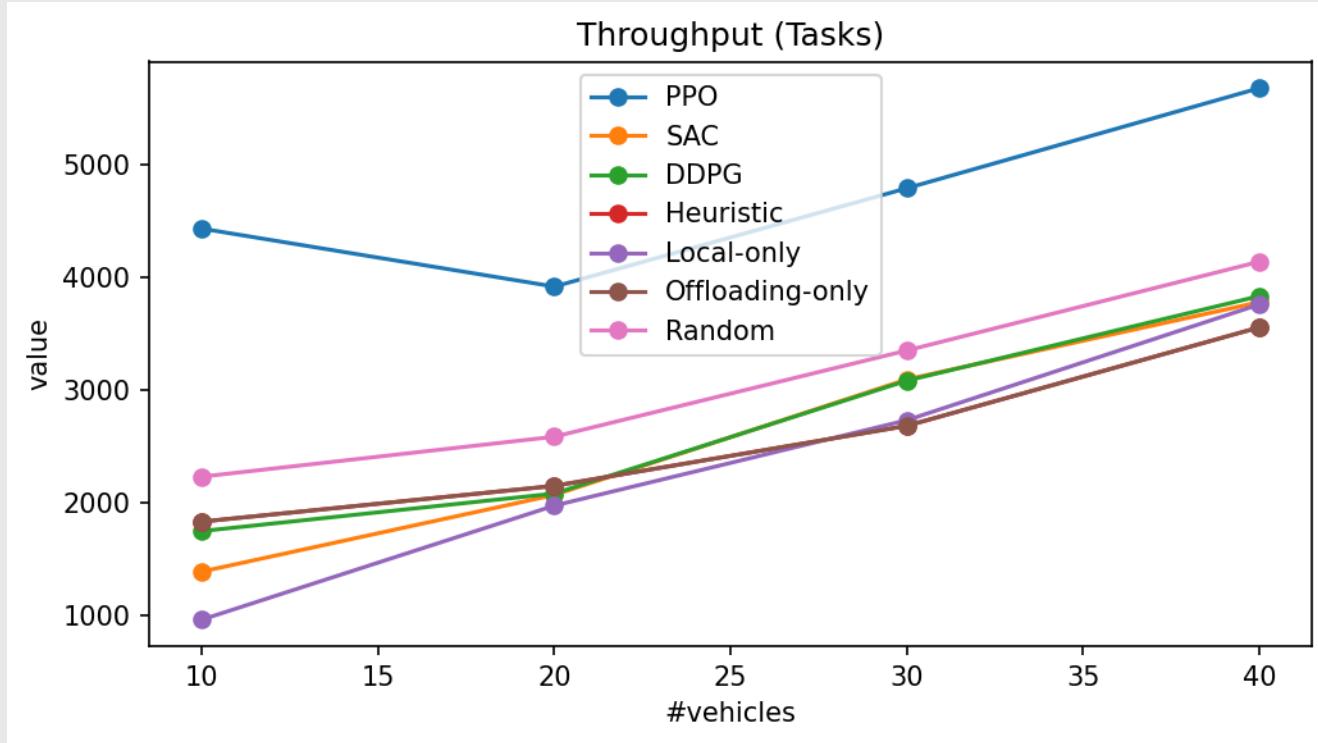
# Energy per Task Comparison



# Deadline Miss Rate Comparison



# Throughput Comparison



# Scalability vs #Vehicles

- **Experiment:**  $V \in \{10, 20, 30, 40, 50\}$ , seeds =  $\{0,1,2\}$ , 5 eval episodes each.
- **Trend highlights:**
  - As **vehicles**↑, **offloading-only** degrades (shared V2I), **local-only** saturates on compute
  - **PPO** maintains **lower miss rate** by balancing local/V2V/edge/cloud
  - Throughput grows with  $V$
- **Interpretation:** Learning when **not** to offload is as important as offloading

# Insights

- **Distance-based power (default) vs learned power (ablation):**
  - Removing the power head shrinks the action space → faster, stabler PPO
- **No-V2V ablation:**
  - Increases queueing at edges/cloud
- **Heuristic (EDF + min-ETA):**
  - Strong at light load; under bursty arrivals and higher V, it misjudges congestion vs PPO
- **Reward normalization (running z-scores):**
  - Critical for stable training on CPU (prevents advantage blow-ups)
- **Takeaway:** The **offload+priority** structure is the main win

# Reproducibility & Engineering

- **One-command pipelines**
  - **Train:** scripts/train.py, train\_sac.py, train\_ddpg.py
  - **Evaluate & metrics:** scripts/eval\_metrics.py
  - **Baselines:** scripts/baseline\_heuristic.py --mode {heuristic,local,offload,random}
  - **Grids & plots:** scripts/run\_grid.py, scripts/plot\_metrics.py, scripts/plot\_vs\_vehicles.py
- **Configs & seeds**
  - YAML configs (Table-III aligned); **config snapshot** saved in each logdir.
  - Multi-seed runs ( $\geq 3$ ) aggregated to mean $\pm$ std tables/plots.
- **Environment**
  - Gymnasium env, Dict observations



**Any Questions?**

# References

1. P. Li, Z. Xiao, X. Wang, K. Huang, Y. Huang and H. Gao, "EPtask: Deep Reinforcement Learning Based Energy-Efficient and Priority-Aware Task Scheduling for Dynamic Vehicular Edge Computing," in IEEE Transactions on Intelligent Vehicles, vol. 9, no. 1, pp. 1830-1846, Jan. 2024, doi: 10.1109/TIV.2023.3321679.